

ST310 Final R Markdown

16923, 17915, 22220

4/6/2021

ST310 Machine Learning Project: Understanding and Predicting Stroke Occurrences in Imbalanced Data

Load the dataset

```
data <- read.csv("healthcare-dataset-stroke-data.csv", header=TRUE)
data <- subset(data, select= -c(1)) #remove the id from the dataframe
```

Load the required packages

```
library(tidyverse); library(ggplot2); library(GGally); library(corrplot); library(SmartEDA)
library(dplyr); library(DataExplorer); library(tibble); library(naniar); library(gridExtra);
library(caret); library(caTools); library(xgboost); library(broom); library(kernlab);
library(MASS); library(modelr); library(glmnet); library(selectiveInference); library(imbalance)
```

EDA

Inspect the data

```
dim(data) # 5110 rows and 11 columns
```

```
## [1] 5110  11
```

```
str(data)
```

```
## 'data.frame':  5110 obs. of  11 variables:
## $ gender      : Factor w/ 3 levels "Female","Male",...: 2 1 2 1 1 2 2 1 1 1 ...
## $ age         : num  67 61 80 49 79 81 74 69 59 78 ...
## $ hypertension : int   0 0 0 0 1 0 1 0 0 0 ...
## $ heart_disease : int   1 0 1 0 0 0 1 0 0 0 ...
## $ ever_married  : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 1 2 2 ...
## $ work_type     : Factor w/ 5 levels "children","Govt_job",...: 4 5 4 4 5 4 4 4 4 ...
## $ Residence_type : Factor w/ 2 levels "Rural","Urban": 2 1 1 2 1 2 1 2 1 2 ...
## $ avg_glucose_level: num   229 202 106 171 174 ...
```

```
## $ bmi : Factor w/ 419 levels "10.3","11.3",...: 240 419 199 218 114 164 148 102 419 116
## $ smoking_status : Factor w/ 4 levels "formerly smoked",...: 1 2 2 3 2 1 2 2 4 4 ...
## $ stroke : int 1 1 1 1 1 1 1 1 1 1 ...
```

```
summary(data)
```

```
##      gender      age      hypertension      heart_disease      ever_married
## Female:2994  Min.   : 0.08  Min.   :0.00000  Min.   :0.00000  No :1757
## Male  :2115  1st Qu.:25.00  1st Qu.:0.00000  1st Qu.:0.00000  Yes:3353
## Other :    1  Median :45.00  Median :0.00000  Median :0.00000
##      Mean   :43.23  Mean   :0.09746  Mean   :0.05401
##      3rd Qu.:61.00  3rd Qu.:0.00000  3rd Qu.:0.00000
##      Max.   :82.00  Max.   :1.00000  Max.   :1.00000
##
##      work_type      Residence_type avg_glucose_level      bmi
## children      : 687  Rural:2514  Min.   : 55.12  N/A   : 201
## Govt_job      : 657  Urban:2596  1st Qu.: 77.25  28.7   : 41
## Never_worked  :  22                Median : 91.89  28.4   : 38
## Private       :2925                Mean   :106.15  26.1   : 37
## Self-employed: 819                3rd Qu.:114.09  26.7   : 37
##      Max.   :271.74  27.6   : 37
##      (Other):4719
##
##      smoking_status      stroke
## formerly smoked: 885  Min.   :0.00000
## never smoked   :1892  1st Qu.:0.00000
## smokes         : 789  Median :0.00000
## Unknown        :1544  Mean   :0.04873
##      3rd Qu.:0.00000
##      Max.   :1.00000
##
```

- Several categorical predictors are wrongly coded as numerical and vice versa.
- Gender has a single datapoint that falls into the level titled 'Other'
- bmi has 201 N/A values
- smoking_status has a category titled 'Unknown'
- Minimum age is 0.08

```
sum(data$age<1) ### 43 people are less than a year old
```

```
## [1] 43
```

Data preparation

```
# Remove the datapoint that falls under the level titled 'Other' for gender
data[data$gender=='Other',] # Identify the row corresponding to this datapoint
```

```
##      gender age hypertension heart_disease ever_married work_type
## 3117 Other  26              0              0          No Private
##      Residence_type avg_glucose_level bmi smoking_status stroke
## 3117      Rural      143.33 22.4 formerly smoked      0
```

```

data <- data[-3117,] # Remove the datapoint
data$gender <- fct_drop(data$gender) # Remove the level 'Other'

# Convert hypertension, heart_disease and stroke to categorical
data$hypertension <- as.factor(data$hypertension)
data$heart_disease <- as.factor(data$heart_disease)
data$stroke <- as.factor(data$stroke)

# Convert bmi too numeric
data$bmi <- as.character(data$bmi)
data$bmi <- as.numeric(data$bmi)

summary(data) # Changes have been made

```

```

##      gender      age      hypertension heart_disease ever_married
## Female:2994   Min.    : 0.08    0:4611          0:4833          No :1756
## Male  :2115   1st Qu.:25.00    1: 498          1: 276          Yes:3353
##                                     Median :45.00
##                                     Mean    :43.23
##                                     3rd Qu.:61.00
##                                     Max.    :82.00
##
##      work_type      Residence_type avg_glucose_level      bmi
## children      : 687   Rural:2513   Min.    : 55.12   Min.    :10.30
## Govt_job      : 657   Urban:2596   1st Qu.: 77.24   1st Qu.:23.50
## Never_worked  :  22               Median : 91.88   Median :28.10
## Private       :2924               Mean    :106.14   Mean    :28.89
## Self-employed: 819               3rd Qu.:114.09   3rd Qu.:33.10
##                                     Max.    :271.74   Max.    :97.60
##                                     NA's     :201
##
##      smoking_status stroke
## formerly smoked: 884   0:4860
## never smoked   :1892   1: 249
## smokes         : 789
## Unknown        :1544
##
##
##

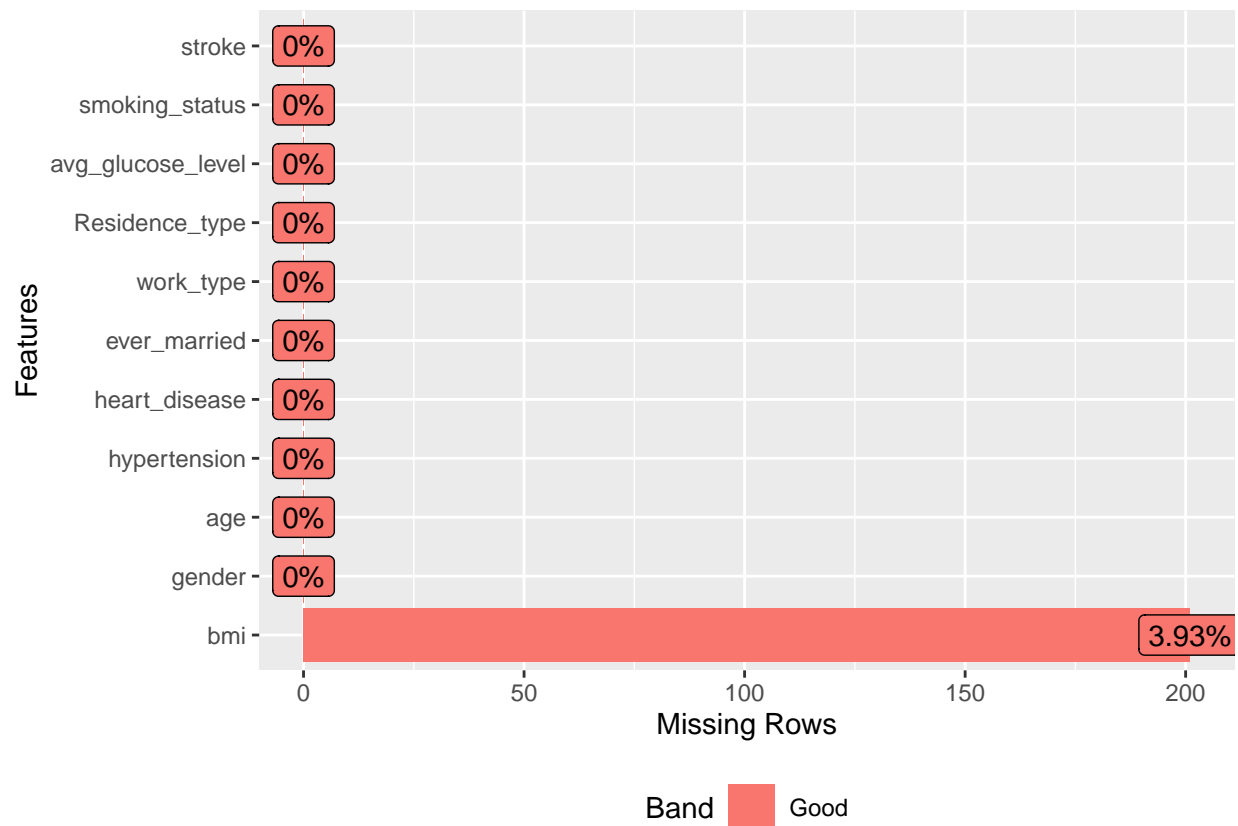
```

Identifying N/A values

```

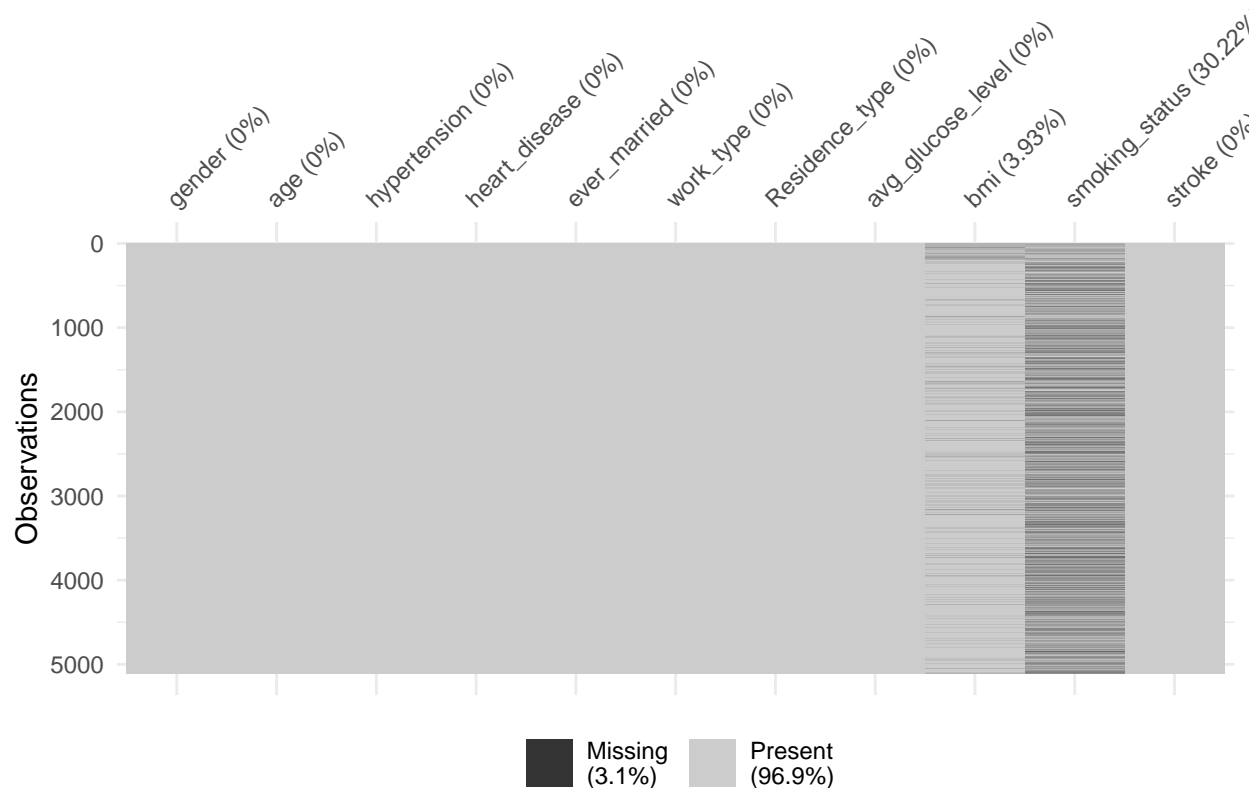
# Plot the amount of missing values for each of the parameters in the dataset
plot_missing(data) # 4% of BMI is missing

```



```
# Replace the "N/A" in bmi & "Unknown" in Smoking status using the naniar package
data_clean <- replace_with_na(data = data, replace = list(bmi = c("N/A"), smoking_status = c("Unknown")))

vis_miss(data_clean) #plot where missing values for each of the parameter in the dataset are + percent
```



#30% of the data is missing for smoking and 4% is missing for bmi

```
sapply(data_clean, function(x) sum(is.na(x)))
```

```
##      gender      age      hypertension      heart_disease
##      0          0          0              0
##  ever_married    work_type  Residence_type avg_glucose_level
##      0          0          0              0
##      bmi      smoking_status      stroke
##      201      1544          0
```

#201 missing values for BMI and 1544 missing values for smoking_status

```
summary(data_clean) # The changes have been made
```

```
##      gender      age      hypertension      heart_disease      ever_married
##  Female:2994  Min.   : 0.08  0:4611      0:4833      No :1756
##  Male   :2115  1st Qu.:25.00  1: 498      1: 276      Yes:3353
##                                     Median :45.00
##                                     Mean   :43.23
##                                     3rd Qu.:61.00
##                                     Max.   :82.00
##
##      work_type      Residence_type      avg_glucose_level      bmi
##  children      : 687      Rural:2513      Min.   : 55.12      Min.   :10.30
```

```
## Govt_job      : 657   Urban:2596      1st Qu.: 77.24      1st Qu.:23.50
## Never_worked : 22                      Median : 91.88      Median :28.10
## Private       :2924                      Mean   :106.14     Mean   :28.89
## Self-employed: 819                      3rd Qu.:114.09    3rd Qu.:33.10
##                                     Max.   :271.74     Max.   :97.60
##                                     NA's   :201
##
##      smoking_status stroke
## formerly smoked: 884   0:4860
## never smoked   :1892   1: 249
## smokes         : 789
## Unknown        :    0
## NA's           :1544
##
##
```

We have 2 possible approaches that we can take when dealing with the N/A values: 1. Remove the rows containing N/A values
2. Impute the bmi/smoking_status values with the most common value (mode)

```
data_remove <- data_clean[complete.cases(data_clean), ]
data_remove$smoking_status <- fct_drop(data_remove$smoking_status) # Remove 'Unknown' as a level of the
dim(data_remove) #3425 observations
```

Method 1 - Remove the rows containing N/A values

```
## [1] 3425 11
```

```
summary(data_remove)
```

```
##      gender      age      hypertension heart_disease ever_married
## Female:2086   Min.   :10.00   0:3017      0:3219      No : 826
## Male  :1339   1st Qu.:34.00   1: 408      1: 206      Yes:2599
##                                     Median :50.00
##                                     Mean   :48.65
##                                     3rd Qu.:63.00
##                                     Max.   :82.00
##      work_type  Residence_type avg_glucose_level      bmi
## children      : 68   Rural:1680   Min.   : 55.12   Min.   :11.50
## Govt_job       : 514   Urban:1745   1st Qu.: 77.23   1st Qu.:25.30
## Never_worked  : 14                      Median : 92.35   Median :29.10
## Private       :2200                      Mean   :108.31   Mean   :30.29
## Self-employed: 629                      3rd Qu.:116.20   3rd Qu.:34.10
##                                     Max.   :271.74   Max.   :92.00
##
##      smoking_status stroke
## formerly smoked: 836   0:3245
## never smoked   :1852   1: 180
## smokes         : 737
##
##
##
```

```
data_impute <- impute_median_at(data_clean, .vars=c("bmi")) #impute bmi at median value using the nania
data_impute <- fill(data_clean, smoking_status) #fills missing values using the previous entry, assumpt
dim(data_impute) #5110
```

Method 2 - Impue values for the N/A values

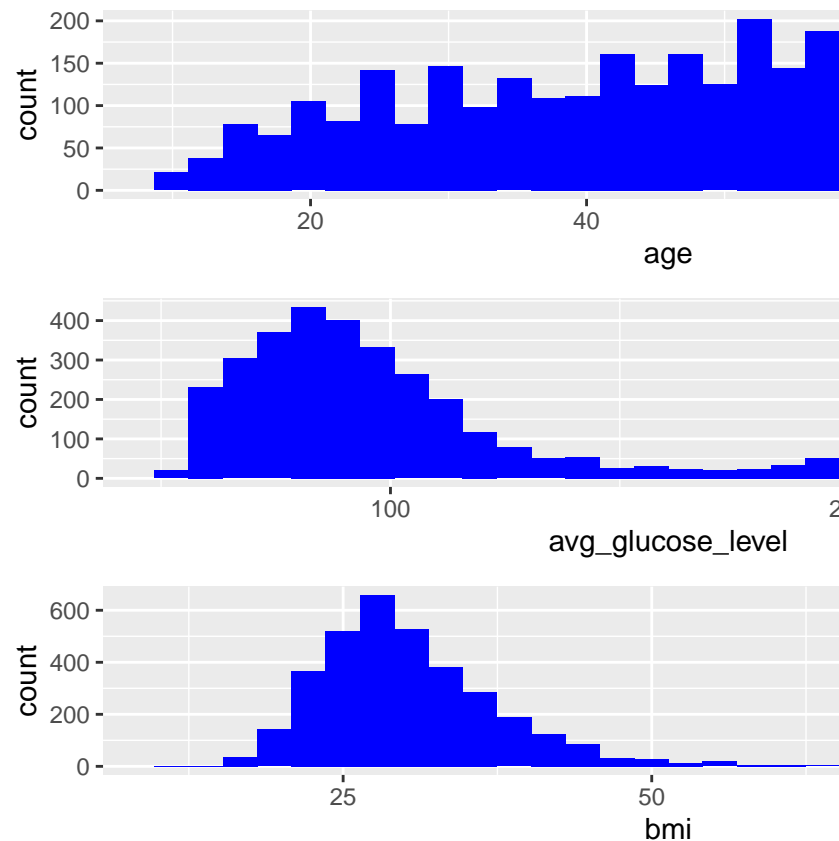
```
## [1] 5109 11
```

We decided to remove the rows that contained N/A values (method 1) for a number of reasons explained in the report

Conducting a univariate analysis of the variables

```
age <- ggplot(data_remove, aes(x=age)) + geom_histogram(fill = "blue")
#approx normally distributed
avg_glucose_level <- ggplot(data_remove, aes(x=avg_glucose_level)) + geom_histogram(fill = "blue")
#bi-modal
bmi <- ggplot(data_remove, aes(x=bmi)) + geom_histogram(fill = "blue")
#approx positively skewed

grid.arrange( age, avg_glucose_level, bmi, ncol=1)
```



Plot histograms for the continuous variables

```
stroke <- ggplot(data_remove, aes(x=stroke)) + geom_bar(stat='count', fill = "red")
gender <- ggplot(data_remove, aes(x=gender)) + geom_bar(stat='count', fill = "red")
hypertension <- ggplot(data_remove, aes(x=hypertension)) + geom_bar(stat='count', fill = "red")
heart_disease <- ggplot(data_remove, aes(x=heart_disease)) + geom_bar(stat='count', fill = "red")
ever_married <- ggplot(data_remove, aes(x=ever_married)) + geom_bar(stat='count', fill = "red")
work_type <- ggplot(data_remove, aes(x=work_type)) + geom_bar(stat='count', fill = "red")
Residence_type <- ggplot(data_remove, aes(x=Residence_type)) + geom_bar(stat='count', fill = "red")
smoking_status <- ggplot(data_remove, aes(x=smoking_status)) + geom_bar(stat='count', fill = "red")

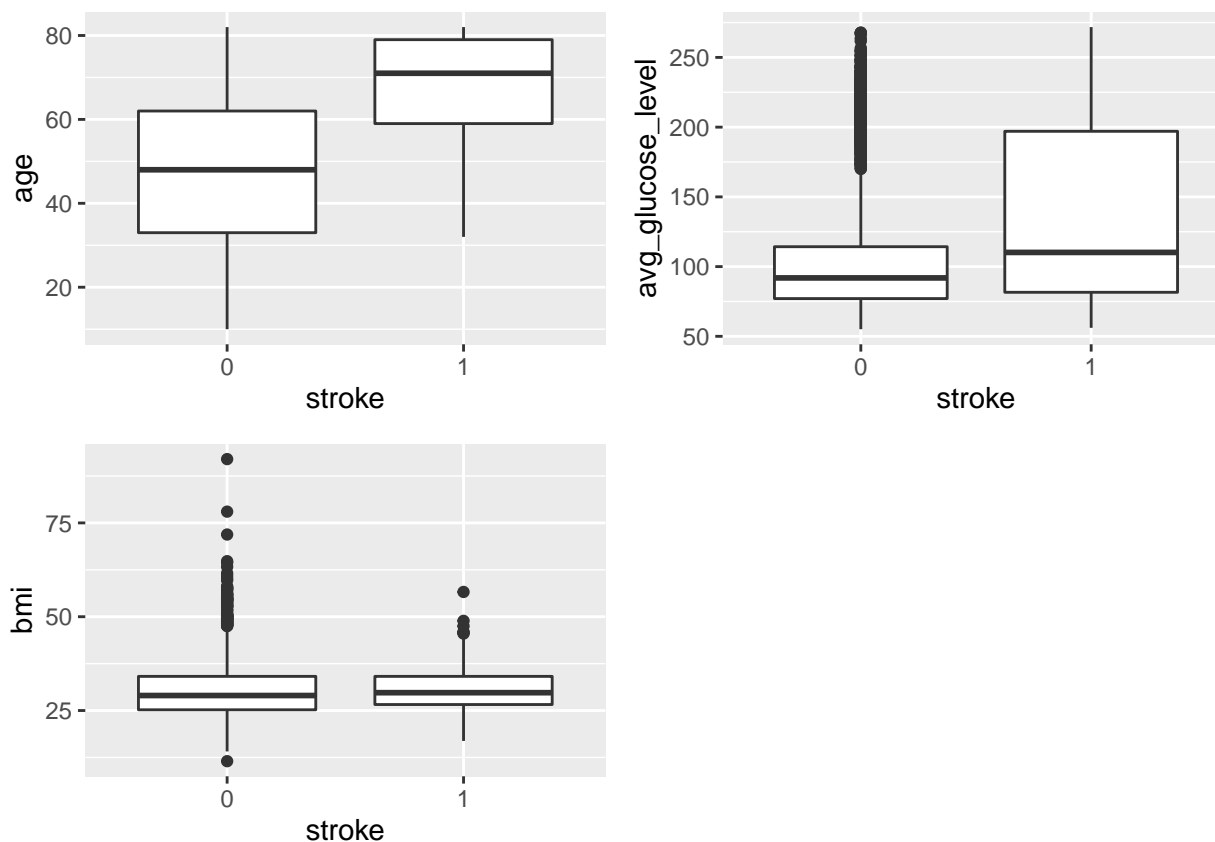
grid.arrange(stroke, gender, hypertension, heart_disease, ever_married, work_type, Residence_type, smoking_status)
```



Plot bar charts to observe the categorical variables

```
age_boxplot <- ggplot(data = data_remove, aes(stroke, age)) +
  geom_boxplot()
avg_glucose_level_boxplot <- ggplot(data = data_remove, aes(stroke, avg_glucose_level)) +
  geom_boxplot()
bmi_boxplot <- ggplot(data = data_remove, aes(stroke, bmi)) +
  geom_boxplot()
grid.arrange( age_boxplot, avg_glucose_level_boxplot, bmi_boxplot, nrow=2)
```


Plot boxplots to determine the relationship between stroke and the continuous predictors

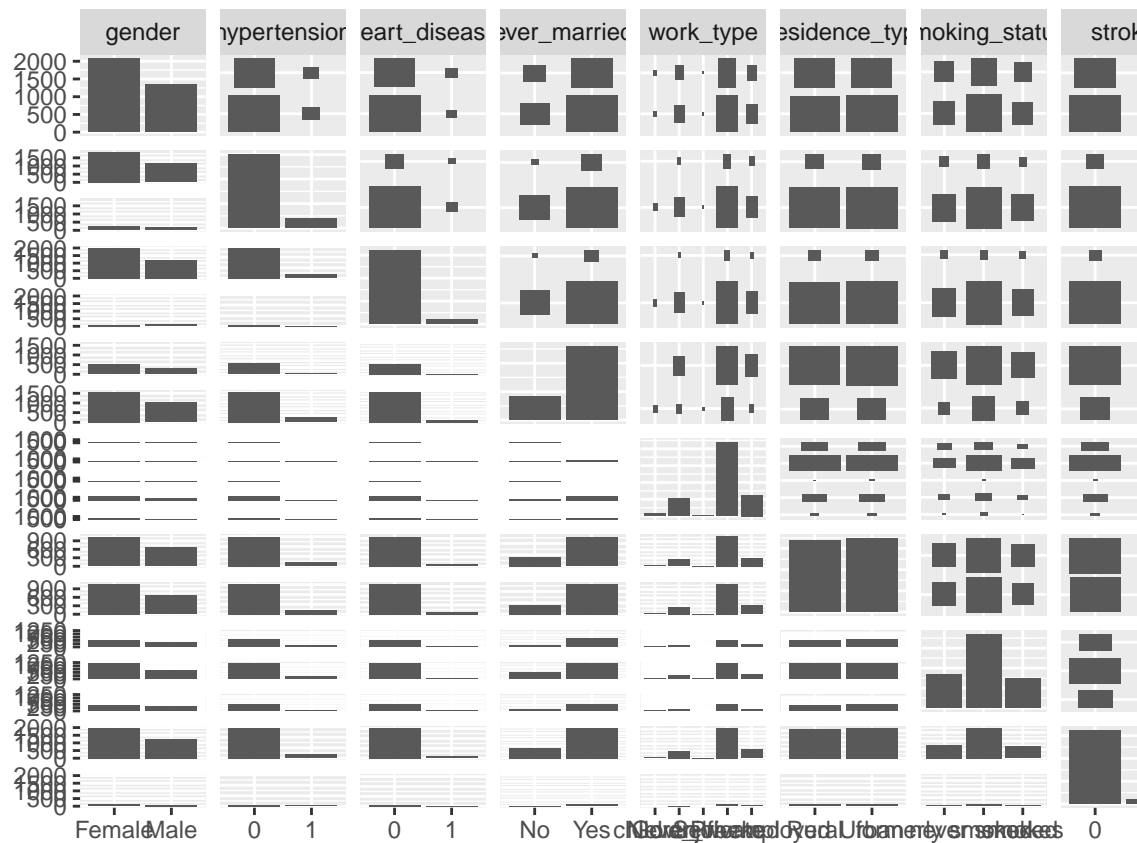


Observe the correlation among variables

Create subsets of the continuous & discrete variables

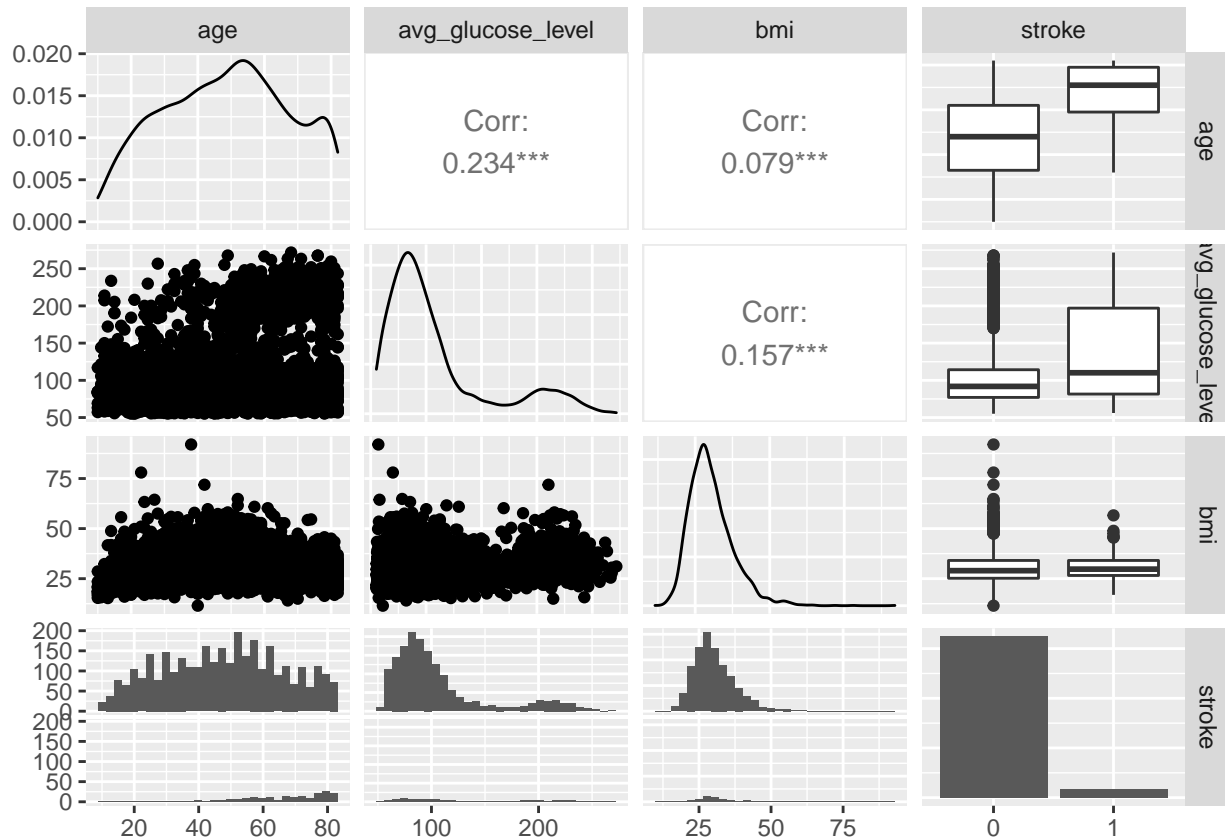
```
contVars <- subset(data_remove, select=c(age, avg_glucose_level, bmi, stroke))
discVars <- subset(data_remove, select=c(gender, hypertension, heart_disease, ever_married, work_type, l
```

```
disccorr <- ggpairs(discVars)
disccorr
```



Pairwise correlations

```
contcorr <- ggpairs(contVars)
contcorr
```



Observe the distribution of people who had a stroke

```
table(data_remove$stroke) #180 total people with a stroke, 3246 with no stroke
```

```
##
##      0      1
## 3245  180
```

```
3245/(180+3245) # Percentage of people who had no stroke
```

```
## [1] 0.9474453
```

- There is a massive class imbalance
- This can be rectified using several methods such as undersampling, oversampling etc.

Dummy model: If we build a model to continuously predict that an individual does not have a stroke, the misclassification rate would be 5% (1-0.9474606)

There are several ways of dealing with the imbalance in the dataset. We will focus on 2 of them 1. Under-sampling - Downsampling the larger class 2. Oversampling - Oversampling the minority class

We will develop our models under both methods

1. Using undersampling

Create a balanced dataset with the same number of observations in both classes (in Stroke) using undersampling

```
stroke_No <- data_remove %>%  
  filter(stroke == 0) %>%  
  sample_n(size = 180)
```

```
stroke_Yes <- data_remove %>%  
  filter(stroke == 1)
```

```
data_under <- rbind(stroke_No, stroke_Yes)  
summary(data_under)
```

```
##      gender      age      hypertension heart_disease ever_married  
## Female:216  Min.   :12.00    0:286          0:317          No : 64  
## Male   :144  1st Qu.:43.00    1: 74          1: 43          Yes:296  
##                               Median :59.00  
##                               Mean   :56.78  
##                               3rd Qu.:73.25  
##                               Max.   :82.00  
##      work_type  Residence_type avg_glucose_level      bmi  
## children      : 4    Rural:172      Min.   : 55.46  Min.   :16.90  
## Govt_job      : 43   Urban:188      1st Qu.: 79.08  1st Qu.:25.50  
## Never_worked  : 1                               Median : 97.66  Median :29.40  
## Private       :226                               Mean   :120.81  Mean   :30.23  
## Self-employed: 86      3rd Qu.:168.29  3rd Qu.:34.10  
##                               Max.   :271.74  Max.   :56.60  
##      smoking_status stroke  
## formerly smoked:105    0:180  
## never smoked   :189    1:180  
## smokes         : 66  
##  
##  
##
```

Create the training and testing data for the dataset obtained using undersampling

```
set.seed(4)  
  
data_under$id <- 1:nrow(data_under) # Create an id  
training_under <- data_under %>% sample_frac(.7)  
testing_under <- anti_join(data_under, training_under, by = 'id')  
training_under <- training_under %>% dplyr::select(-id)  
testing_under <- testing_under %>% dplyr::select(-id)  
beta_testing_under <- testing_under %>% dplyr::select(-stroke)  
  
#summary(training_under)  
#dim(training_under)  
#summary(testing_under)  
#dim(testing_under)
```

2. Using oversampling

Create a balanced dataset with the same number of observations in both classes using oversampling

```
data_over <- data_remove

# Compute the imbalance ratio of stroke
imbalanceRatio(as.data.frame(data_over), classAttr = "stroke")

## [1] 0.05546995

# Name the levels of stroke
data_over$stroke <- as.factor(ifelse(data_over$stroke == 0, "no", "yes"))

# Put variables as correct format
data_over$gender <- as.factor(data_over$gender)
data_over$hypertension <- as.factor(data_over$hypertension)
data_over$heart_disease <- as.factor(data_over$heart_disease)
data_over$ever_married <- as.factor(data_over$ever_married)
data_over$work_type <- as.factor(data_over$work_type)
data_over$Residence_type <- as.factor(data_over$Residence_type)
data_over$smoking_status <- as.factor(data_over$smoking_status)
data_over <- as.data.frame(lapply(data_over, as.numeric))

data_over <- oversample(as.data.frame(data_over), classAttr = "stroke", ratio = 1, method = "MWMOTE")

data_over$stroke <- as.factor(data_over$stroke)

table(data_over$stroke)

##
##      1      2
## 3245 3245

#Bot as is no shown, there are now equally as many stroke cases as non-stroke cases
summary(data_over)
```

```
##      gender      age      hypertension      heart_disease
## Min.   :1.000   Min.   :10.00   Min.   :1.000   Min.   :1.000
## 1st Qu.:1.000   1st Qu.:47.00   1st Qu.:1.000   1st Qu.:1.000
## Median :1.072   Median :63.24   Median :1.000   Median :1.000
## Mean   :1.408   Mean   :58.96   Mean   :1.208   Mean   :1.144
## 3rd Qu.:2.000   3rd Qu.:75.00   3rd Qu.:1.292   3rd Qu.:1.000
## Max.   :2.000   Max.   :82.00   Max.   :2.000   Max.   :2.000
## ever_married  work_type  Residence_type  avg_glucose_level
## Min.   :1.000   Min.   :1.000   Min.   :1.000   Min.   : 55.12
## 1st Qu.:2.000   1st Qu.:4.000   1st Qu.:1.000   1st Qu.: 80.88
## Median :2.000   Median :4.000   Median :1.547   Median : 99.72
## Mean   :1.827   Mean   :3.928   Mean   :1.512   Mean   :124.18
## 3rd Qu.:2.000   3rd Qu.:4.340   3rd Qu.:2.000   3rd Qu.:180.95
## Max.   :2.000   Max.   :5.000   Max.   :2.000   Max.   :271.74
```

```
##      bmi      smoking_status  stroke
##  Min.   :11.50   Min.       :1.000   1:3245
##  1st Qu.:26.13   1st Qu.:1.478   2:3245
##  Median :29.30   Median :2.000
##  Mean   :30.21   Mean    :1.944
##  3rd Qu.:33.23   3rd Qu.:2.032
##  Max.   :92.00   Max.     :3.000
```

Create the training and testing data for the dataset obtained using oversampling

```
data_over$id <- 1:nrow(data_over)
training_over <- data_over %>% sample_frac(.7)
testing_over <- anti_join(data_over, training_over, by = 'id')
training_over <- training_over %>% dplyr::select(-id)
testing_over <- testing_over %>% dplyr::select(-id)
beta_testing_over <- testing_over %>% dplyr::select(-stroke)

dim(training_over) #4543 observations
```

```
## [1] 4543  11
```

```
dim(testing_over) #1947 observations
```

```
## [1] 1947  11
```

Models developed using undersampling

Model 1: Logistic regression

stepAIC in the MASS package was used to obtain the model that contains the most contributive predictors by minimising AIC

```
# Build a glm with all the predictors
glm_all_under <- glm(stroke~., family=binomial(link = "logit"), data=training_under)

# Build a glm with only the most contributive predictors
glm_under <- glm_all_under %>% stepAIC(trace = FALSE)
# trace = FALSE allows the function to provide only the model with the lowest AIC and none of the intermediate models
summary(glm_under)
```

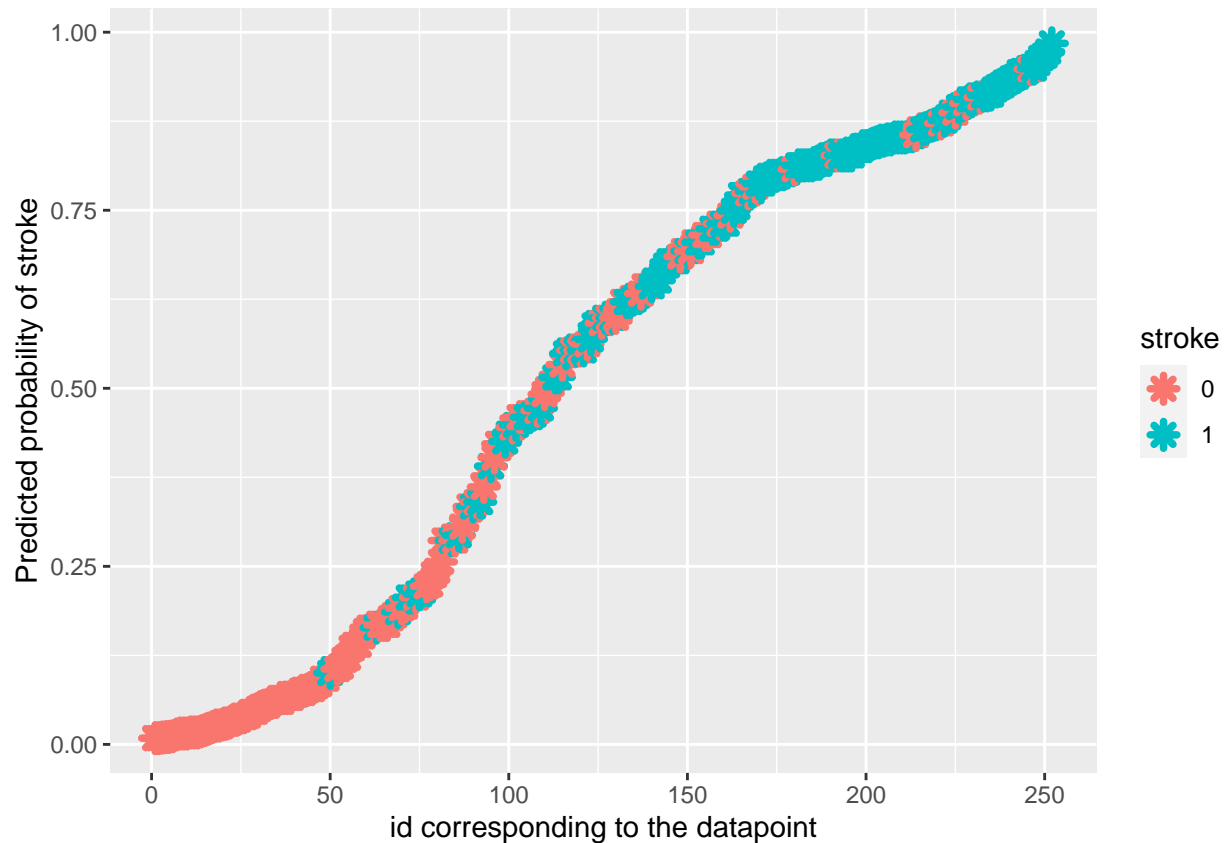
```
##
## Call:
## glm(formula = stroke ~ age + hypertension + smoking_status, family = binomial(link = "logit"),
##      data = training_under)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4330  -0.6003   0.2654   0.6639   2.1441
##
## Coefficients:
```

```
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -5.70361    0.81788  -6.974 3.09e-12 ***
## age               0.09392    0.01224   7.674 1.67e-14 ***
## hypertension1     0.85917    0.42102   2.041 0.04129 *
## smoking_statusnever smoked -0.15206    0.37382  -0.407 0.68418
## smoking_statussmokes 1.37822    0.51234   2.690 0.00714 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 349.20  on 251  degrees of freedom
## Residual deviance: 221.71  on 247  degrees of freedom
## AIC: 231.71
##
## Number of Fisher Scoring iterations: 5
```

```
#Create a visual representation of the logistic model
predicted.data <- data.frame(
  probability.of.stroke = glm_under$fitted.values,
  stroke=training_under$stroke)

predicted.data <- predicted.data[
  order(predicted.data$probability.of.stroke, decreasing=FALSE),]
predicted.data$rank <- 1:nrow(predicted.data)

ggplot(data=predicted.data, aes(x=rank, y=probability.of.stroke)) +
  geom_point(aes(color=stroke), alpha=1, shape=8, size = 2, stroke = 2) +
  xlab("id corresponding to the datapoint") +
  ylab("Predicted probability of stroke")
```



```
# Predict the testing data using the glm
pred_glm_under = predict(glm_under, testing_under, type="response")
pred_glm_under <- ifelse((pred_glm_under>=0.3), 1, 0) # This threshold value was one of the values that

# Create a confusion matrix to assess the performance of the model
conf_matrix_glm_under <- table(pred_glm_under, testing_under$stroke, deparse.level = 2)
conf_matrix_glm_under
```

```
##           testing_under$stroke
## pred_glm_under  0  1
##           0 29  4
##           1 28 47
```

Comparing glm_all_under and glm_under

```
rbind(glance(glm_all_under), glance(glm_under))
```

```
## # A tibble: 2 x 8
##   null.deviance df.null logLik   AIC   BIC deviance df.residual  nobs
##   <dbl>      <int> <dbl> <dbl> <dbl>   <dbl>      <int> <int>
## 1     349.      251 -109.  247.  300.    217.      237  252
## 2     349.      251 -111.  232.  249.    222.      247  252
```


Model 2: Linear model to predict stroke implemented using gradient descent

Preprocess the data

```
x_var_train_for_lm_under <- training_under %>%
  dplyr::select(age, bmi)

# Create a design matrix for the training data
X_train_for_lm_under <- model.matrix(~. -1, data = x_var_train_for_lm_under) # -1 removes the intercept

# Convert stroke to numeric
train_under_2 <- training_under
train_under_2$stroke <- as.numeric(train_under_2$stroke)
Y_train_for_lm_under <- train_under_2 %>% pull(stroke)

x_var_test_for_lm_under <- testing_under %>%
  dplyr::select(age, bmi)

# Create a design matrix for the testing data
X_test_for_lm_under <- model.matrix(~. -1, data = x_var_test_for_lm_under)

# Convert stroke to numeric
test_under_2 <- testing_under
test_under_2$stroke <- as.numeric(test_under_2$stroke)

Y_test_for_lm_under <- test_under_2 %>% pull(stroke)

dim(X_train_for_lm_under) #252

## [1] 252  2

dim(X_test_for_lm_under) #108

## [1] 108  2

# Define the least squares function
least_squares_gradient <- function(x, y, beta) {
  -2 * t(x) %*% (y - x %*% beta)
}
# t(x) is the transpose of x

# Define the loss function
least_squares_loss <- function(x, y, beta) {
  sum((y - x %*% beta)^2)
}

# Initialize coefficients
gamma = 0.000001 # this is the step size
p = 2 # This is the number of predictors

beta0 <- rep(0, p) # This is the vector of all 0s
previous_loss <- least_squares_loss(X_train_for_lm_under, Y_train_for_lm_under, beta0) # Loss function
```

```

grad0 <- least_squares_gradient(X_train_for_lm_under, Y_train_for_lm_under, beta0) # Initialise the gra
beta1 <- beta0 - gamma * grad0
next_loss <- least_squares_loss(X_train_for_lm_under, Y_train_for_lm_under, beta1)
previous_beta <- beta1
steps <- 1

while (abs(previous_loss - next_loss) > 0.00001) {
  gradn <- least_squares_gradient(X_train_for_lm_under, Y_train_for_lm_under, previous_beta)
  # Refine update by allowing step size to change at each iteration. Make step size a sequence of numbe
  next_beta <- previous_beta - (0.99)^steps * gradn / sqrt(sum(gradn^2))
  # We rescale the gradient i.e. use sqrt(sum(gradn^2)) to prevent the algorithm from diverging
  steps <- steps + 1
  previous_beta <- next_beta
  previous_loss <- next_loss
  next_loss <- least_squares_loss(X_train_for_lm_under, Y_train_for_lm_under, next_beta)
}

# Predict the expected values for the testing data
pred_lm_grad_desc_under = X_test_for_lm_under %*% previous_beta
pred_lm_grad_desc_under = round(as.integer(pred_lm_grad_desc_under))
pred_lm_grad_desc_under = as.factor(pred_lm_grad_desc_under)

# Create a confusion matrix to assess the performance of the model on the testing data
table(pred_lm_grad_desc_under, testing_under$stroke, deparse.level = 2)

```

```

##               testing_under$stroke
## pred_lm_grad_desc_under  0  1
##               0  9  0
##               1 48 48
##               2  0  3

```

Variable that is rounded to 2, is actually a 1, as 0 and 1 are the only outcomes

Model 3: Lasso (least absolute shrinkage and selection operator) regression

Preprocess this data to use for lasso

```

x_var_train_for_lasso_under <- training_under %>%
  dplyr::select(-stroke)

# Create a design matrix for the training data
X_train_for_lasso_under <- model.matrix(~. -1, data = x_var_train_for_lasso_under) # -1 removes the int

Y_train_for_lasso_under <- training_under %>% pull(stroke)

x_var_test_for_lasso_under <- testing_under %>%
  dplyr::select(-stroke)

# Create a design matrix for the testing data
X_test_for_lasso_under <- model.matrix(~. -1, data = x_var_test_for_lasso_under) #

Y_test_for_lasso_under <- testing_under %>% pull(stroke)

```

```
dim(X_train_for_lasso_under) #252
```

```
## [1] 252 15
```

```
dim(X_test_for_lasso_under) #108
```

```
## [1] 108 15
```

The model was created using lambda.min as the best lambda

```
#Perform 10 fold cross validation using the misclassification rate to find the best lambda
lasso_cv_under = cv.glmnet(X_train_for_lasso_under,
                           Y_train_for_lasso_under,
                           family = "binomial",
                           type.measure = "class") # type.measure = "class" allows us to use the misclassification rate
plot(lasso_cv_under)
```

```
# Fit the lasso model on the training data
lasso_under <- glmnet(X_train_for_lasso_under,
                     Y_train_for_lasso_under,
                     alpha = 1,
                     family = "binomial",
                     lambda = lasso_cv_under$lambda.min)
```

```
# Predict the testing data using the glm
pred_lasso_under <- lasso_under %>% predict(newx = X_test_for_lasso_under)
pred_lasso_under <- ifelse(pred_lasso_under >= 0.3, 1, 0)
```

```
# Create a confusion matrix to assess the performance of the model
table(pred_lasso_under, Y_test_for_lasso_under, deparse.level = 2)
```

```
##               Y_test_for_lasso_under
## pred_lasso_under  0  1
##               0 42 14
##               1 15 37
```

Model 4: kernel method

```
# Build the model using training data
ksvm_under <- ksvm(stroke ~ age + hypertension + heart_disease + avg_glucose_level + smoking_status , k
ksvm_under
```

```
## Support Vector Machine object of class "ksvm"
##
## SV type: C-svc (classification)
## parameter : cost C = 1
##
## Gaussian Radial Basis kernel function.
## Hyperparameter : sigma = 0.483051818661611
```

```
##
## Number of Support Vectors : 146
##
## Objective Function Value : -111.6252
## Training error : 0.178571

# Predict the testing data using the model built
pred_kvsm_under <- predict(kvsm_under, type = 'response', newdata = testing_under)

# Create a confusion matrix to assess the performance of the model
conf_matrix_kvsm_under<-table(Predicted=pred_kvsm_under,Reference=testing_under[,11])
confusionMatrix(conf_matrix_kvsm_under, stroke = 1)

## Confusion Matrix and Statistics
##
##           Reference
## Predicted  0  1
##           0 41  8
##           1 16 43
##
##           Accuracy : 0.7778
##           95% CI : (0.6876, 0.8521)
##       No Information Rate : 0.5278
##       P-Value [Acc > NIR] : 6.833e-08
##
##           Kappa : 0.5578
##
##  McNemar's Test P-Value : 0.153
##
##           Sensitivity : 0.7193
##           Specificity : 0.8431
##           Pos Pred Value : 0.8367
##           Neg Pred Value : 0.7288
##           Prevalence : 0.5278
##           Detection Rate : 0.3796
##       Detection Prevalence : 0.4537
##           Balanced Accuracy : 0.7812
##
##           'Positive' Class : 0
##
```

Model 5: random forest building

```
# Set the parameters for the train function
rftunegrid <- data.frame(
  .mtry=c(2,3,4,5,6), .splitrule="gini", .min.node.size=5
)
rfcontrol <- trainControl(
  method="oob", number=5, verboseIter=TRUE
)
```

```
# Build the model
randomforest_under <- train(
  stroke~., training_under, method="ranger", tuneLength=3, tuneGrid= rftunegrid, trControl=rfcontrol
)
```

```
## + : mtry=2, splitrule=gini, min.node.size=5
## - : mtry=2, splitrule=gini, min.node.size=5
## + : mtry=3, splitrule=gini, min.node.size=5
## - : mtry=3, splitrule=gini, min.node.size=5
## + : mtry=4, splitrule=gini, min.node.size=5
## - : mtry=4, splitrule=gini, min.node.size=5
## + : mtry=5, splitrule=gini, min.node.size=5
## - : mtry=5, splitrule=gini, min.node.size=5
## + : mtry=6, splitrule=gini, min.node.size=5
## - : mtry=6, splitrule=gini, min.node.size=5
## Aggregating results
## Selecting tuning parameters
## Fitting mtry = 3, splitrule = gini, min.node.size = 5 on full training set
```

```
randomforest_under
```

```
## Random Forest
##
## 252 samples
## 10 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling results across tuning parameters:
##
##   mtry  Accuracy  Kappa
##   2     0.7420635  0.4815799
##   3     0.7500000  0.4979127
##   4     0.7420635  0.4821700
##   5     0.7500000  0.4982935
##   6     0.7500000  0.4979127
##
## Tuning parameter 'splitrule' was held constant at a value of gini
##
## Tuning parameter 'min.node.size' was held constant at a value of 5
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were mtry = 3, splitrule = gini
## and min.node.size = 5.
```

```
# Predict the testing data using the model built
randomforest_under_prediction <- predict(randomforest_under, newdata=beta_testing_under)

# Create a confusion matrix to assess the performance of the model on the testing data
confusionMatrix(randomforest_under_prediction, factor(testing_under[["stroke"]]), positive = "1")
```

```
## Confusion Matrix and Statistics
##
```

```

##           Reference
## Prediction  0  1
##           0 40  8
##           1 17 43
##
##           Accuracy : 0.7685
##           95% CI : (0.6775, 0.8443)
##           No Information Rate : 0.5278
##           P-Value [Acc > NIR] : 2.104e-07
##
##           Kappa : 0.5399
##
## Mcnemar's Test P-Value : 0.1096
##
##           Sensitivity : 0.8431
##           Specificity : 0.7018
##           Pos Pred Value : 0.7167
##           Neg Pred Value : 0.8333
##           Prevalence : 0.4722
##           Detection Rate : 0.3981
##           Detection Prevalence : 0.5556
##           Balanced Accuracy : 0.7724
##
##           'Positive' Class : 1
##

```

Model 6: extreme gradient boosting tree

```

# Set the parameters for the train function
xgbgrid <- expand.grid(
  nrounds = 3500, max_depth = 7, eta = 0.01, gamma = 0.01,
  colsample_bytree = 0.75, min_child_weight = 0, subsample = 0.5
)

xgbcontrol <- trainControl(
  method = "cv", number = 5
)

# Build the model
xgb_under <- train(
  stroke ~ ., training_under, method = "xgbTree", tuneLength = 3, tuneGrid = xgbgrid, trControl = xgbcontrol
)

xgb_under

```

```

## eXtreme Gradient Boosting
##
## 252 samples
## 10 predictor
## 2 classes: '0', '1'
##
## No pre-processing

```

```

## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 201, 202, 201, 202, 202
## Resampling results:
##
##   Accuracy   Kappa
##   0.7301176  0.4584377
##
## Tuning parameter 'nrounds' was held constant at a value of 3500
##
## Tuning parameter 'min_child_weight' was held constant at a value of 0
##
## Tuning parameter 'subsample' was held constant at a value of 0.5

# Predict the testing data using the model built
xgb_pred <- predict(xgb_under, newdata = beta_testing_under)

# Create a confusion matrix to assess the performance of the model on the testing data
confusionMatrix(xgb_pred, factor(testing_under[["stroke"]]), positive = "1")

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 38   7
##           1 19  44
##
##               Accuracy : 0.7593
##               95% CI : (0.6675, 0.8363)
##       No Information Rate : 0.5278
##       P-Value [Acc > NIR] : 6.161e-07
##
##               Kappa : 0.5229
##
##  Mcnemar's Test P-Value : 0.03098
##
##       Sensitivity : 0.8627
##       Specificity : 0.6667
##       Pos Pred Value : 0.6984
##       Neg Pred Value : 0.8444
##       Prevalence : 0.4722
##       Detection Rate : 0.4074
##       Detection Prevalence : 0.5833
##       Balanced Accuracy : 0.7647
##
##       'Positive' Class : 1
##

```

Models developed using oversampling

Model 1: Logistic regression

```
glm_all_over <- glm(stroke~., family=binomial(link = "logit"), data=training_over)
glm_over <- glm_all_over %>% stepAIC(trace = FALSE)
summary(glm_over)
```

```
##
## Call:
## glm(formula = stroke ~ gender + age + hypertension + heart_disease +
##     ever_married + work_type + avg_glucose_level + smoking_status,
##     family = binomial(link = "logit"), data = training_over)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7563  -0.6286  -0.1230   0.7218   2.5002
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -8.2111915   0.4315030  -19.029  < 2e-16 ***
## gender        -0.1903558   0.0897963   -2.120  0.03402 *
## age           0.0935111   0.0033173   28.189  < 2e-16 ***
## hypertension   0.5992638   0.1088342    5.506 3.67e-08 ***
## heart_disease  0.7698440   0.1366198    5.635 1.75e-08 ***
## ever_married  -0.3236938   0.1485721   -2.179  0.02935 *
## work_type      0.1309077   0.0471366    2.777  0.00548 **
## avg_glucose_level 0.0066687  0.0007319    9.111  < 2e-16 ***
## smoking_status  0.1918944   0.0610377    3.144  0.00167 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6296.6  on 4542  degrees of freedom
## Residual deviance: 4076.0  on 4534  degrees of freedom
## AIC: 4094
##
## Number of Fisher Scoring iterations: 5
```

```
# create a confusion matrix
pred_glm_over = predict(glm_over, testing_over, type="response")
pred_glm_over <- ifelse((pred_glm_over>=0.3), 1, 0) # This value was preset
conf_matrix_over <- table(pred_glm_over, testing_over$stroke, deparse.level = 2)
conf_matrix_over
```

```
##              testing_over$stroke
## pred_glm_over    1    2
##                0 621  63
##                1 314 949
```

Model 2: Linear regression to predict using gradient descent

Preprocess the data


```

x_var_train_for_lm_over <- training_over %>%
  dplyr::select(age, bmi)

# Create a design matrix for the training data
X_train_for_lm_over <- model.matrix(~. -1, data = x_var_train_for_lm_over) # -1 removes the intercept

train_over_2 <- training_over
train_over_2$stroke <- as.numeric(train_over_2$stroke)

Y_train_for_lm_over <- train_over_2 %>% pull(stroke)

x_var_test_for_lm_over <- testing_over %>%
  dplyr::select(age, bmi)

# Create a design matrix for the testing data
X_test_for_lm_over <- model.matrix(~. -1, data = x_var_test_for_lm_over) #

test_over_2 <- training_over
test_over_2$stroke <- as.numeric(test_over_2$stroke)

Y_test_for_lm_over <- test_over_2 %>% pull(stroke)

dim(X_train_for_lm_over) #4543

## [1] 4543    2

dim(X_test_for_lm_over) #1947

## [1] 1947    2

# Initialize coefficients
gamma = 0.000001 # this is the step size
p = 2 # This is the number of predictors

beta0 <- rep(0, p) # This is the vector of all 0s
previous_loss <- least_squares_loss(X_train_for_lm_over, Y_train_for_lm_over, beta0) # Loss function at
grad0 <- least_squares_gradient(X_train_for_lm_over, Y_train_for_lm_over, beta0) # Initialise gradient
beta1 <- beta0 - gamma * grad0
next_loss <- least_squares_loss(X_train_for_lm_over, Y_train_for_lm_over, beta1)
previous_beta <- beta1
steps <- 1

while (abs(previous_loss - next_loss) > 0.00001) {
  gradn <- least_squares_gradient(X_train_for_lm_over, Y_train_for_lm_over, previous_beta)
  # Refine update by allowing step size to change at each iteration. Make step size a sequence of numbers
  next_beta <- previous_beta - (0.99)^steps * gradn / sqrt(sum(gradn^2))
  # We rescale the gradient i.e. use sqrt(sum(gradn^2)) to prevent the algorithm from diverging
  steps <- steps + 1
  previous_beta <- next_beta
  previous_loss <- next_loss
  next_loss <- least_squares_loss(X_train_for_lm_over, Y_train_for_lm_over, next_beta)
}

```

```
# Use the model the parameter estimates obtained using gradient descent to predict the testing data
pred_lm_grad_desc_over = X_test_for_lm_over %*% previous_beta
pred_lm_grad_desc_over = round(as.integer(pred_lm_grad_desc_over))
pred_lm_grad_desc_over = as.factor(pred_lm_grad_desc_over + 1)
```

```
# Create a confusion matrix to see how well the testing data is predicted
table(pred_lm_grad_desc_over, testing_over$stroke, deparse.level = 2)
```

```
##               testing_over$stroke
## pred_lm_grad_desc_over    1     2
##               1  264    3
##               2  671 1008
##               3    0    1
```

Model 3: Lasso (least absolute shrinkage and selection operator) regression

Preprocess this data to use for lasso

```
x_var_train_for_lasso_over <- training_over %>%
  dplyr::select(-stroke)
# Create the design matrix for the training data
X_train_for_lasso_over <- model.matrix(~. -1, data = x_var_train_for_lasso_over)
```

```
Y_train_for_lasso_over <- training_over %>% pull(stroke)
```

```
x_var_test_for_lasso_over <- testing_over %>%
  dplyr::select(-stroke)
# Create the design matrix for the testing data
X_test_for_lasso_over <- model.matrix(~. -1, data = x_var_test_for_lasso_over)
```

```
Y_test_for_lasso_over <- testing_over %>% pull(stroke)
```

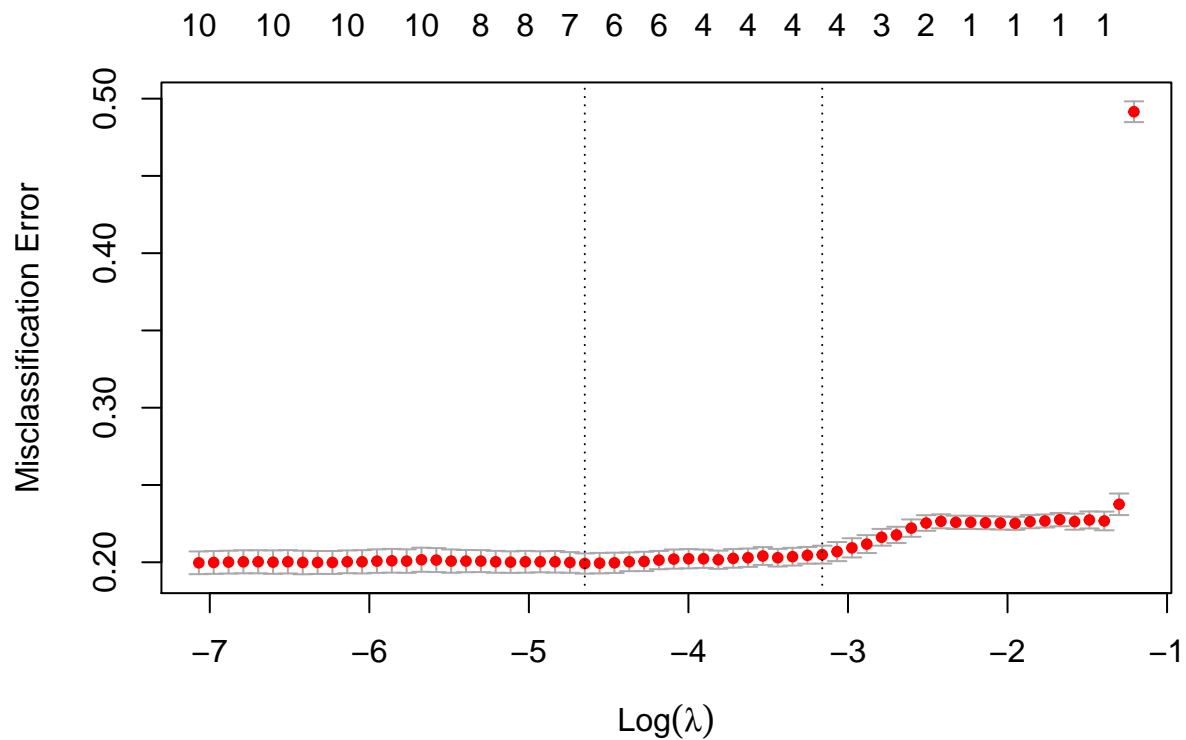
```
dim(X_train_for_lasso_over)
```

```
## [1] 4543    10
```

```
dim(X_test_for_lasso_over)
```

```
## [1] 1947    10
```

```
#Perform 10 fold cross validation using the misclassification rate to find the best lambda
lasso_cv_over = cv.glmnet(X_train_for_lasso_over,
  Y_train_for_lasso_over,
  family = "binomial",
  type.measure = "class") # type.measure = "class" allows us to use the misclassification rate
plot(lasso_cv_over)
```



```
# Fit the lasso model on the training data
lasso_over <- glmnet(X_train_for_lasso_over,
                     Y_train_for_lasso_over,
                     alpha = 1,
                     family = "binomial",
                     lambda = lasso_cv_over$lambda.min)

coef(lasso_over) # Observe what coefficients are included in the model
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept)  -7.586969947
## gender       -0.006107622
## age          0.085735590
## hypertension  0.507606465
## heart_disease 0.593544989
## ever_married  .
## work_type     0.057024201
## Residence_type .
## avg_glucose_level 0.005307559
## bmi          .
## smoking_status 0.072539170
```

```
# Making predictions based off of testing data
pred_lasso_over <- lasso_over %>% predict(newx = X_test_for_lasso_over)
```

```
pred_lasso_over <- ifelse(pred_lasso_over >= 0.3, 1, 0)
table(pred_lasso_over, Y_test_for_lasso_over, deparse.level = 2)
```

```
##           Y_test_for_lasso_over
## pred_lasso_over    1    2
##           0 775 201
##           1 160 811
```

Model 4: Kernel Methods

```
# Find the value of gamma that minimizes the number of false negatives
try <- -100:100 # List of possible values for gamma
false_neg <- data.frame()
error <- data.frame()
for (gam in c(-100:100)){
  # Create the model using training data
  kvsm_over <- ksvm(stroke ~ ., kernel = 'rbfdot', data = training_over, gamma = gam)
  pred_kvsm_over <- predict(kvsm_over, type = 'response', newdata = training_over)
  pred_kvsm_over <- round(as.numeric(pred_kvsm_over))
  pred_kvsm_over <- as.factor(pred_kvsm_over)
  conf_mat <- as.factor(training_over[,11])
  # Subtract accuracy from the false positives
  ans = (confusionMatrix(pred_kvsm_over, conf_mat)[[2]][3]) - confusionMatrix(pred_kvsm_over, conf_mat)
  # As accuracy < 1, this will give us the value with the highest accuracy out of those with the joint
  false_neg <- rbind(false_neg, ans)
}

try[which(false_neg == min(false_neg))] # Choose one of these values as gamma
```

```
## [1] -100 -91 38 47
```

```
kvsm_over <- ksvm(stroke ~ ., kernel = 'rbfdot', data = training_over, gamma = 10)
pred_kvsm_over <- predict(kvsm_over, type = 'response', newdata = testing_over)
pred_kvsm_over <- round(as.numeric(pred_kvsm_over))
pred_kvsm_over <- as.factor(pred_kvsm_over)
conf_mat <- as.factor(testing_over[,11])
confusionMatrix(pred_kvsm_over, conf_mat)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    1    2
##           1 836 105
##           2  99 907
##
##           Accuracy : 0.8952
##           95% CI : (0.8808, 0.9085)
##           No Information Rate : 0.5198
##           P-Value [Acc > NIR] : <2e-16
##
```

```
##           Kappa : 0.7902
##
## Mcnemar's Test P-Value : 0.7263
##
##           Sensitivity : 0.8941
##           Specificity : 0.8962
##           Pos Pred Value : 0.8884
##           Neg Pred Value : 0.9016
##           Prevalence : 0.4802
##           Detection Rate : 0.4294
##           Detection Prevalence : 0.4833
##           Balanced Accuracy : 0.8952
##
##           'Positive' Class : 1
##
```

Model 5: Random Forest

```
# Build the model using the training data
randomforest_over <- train(
  stroke~., training_over, method="ranger", tuneLength=3, tuneGrid= rftunegrid, trControl=rfcontrol
)
```

```
## + : mtry=2, splitrule=gini, min.node.size=5
## - : mtry=2, splitrule=gini, min.node.size=5
## + : mtry=3, splitrule=gini, min.node.size=5
## - : mtry=3, splitrule=gini, min.node.size=5
## + : mtry=4, splitrule=gini, min.node.size=5
## - : mtry=4, splitrule=gini, min.node.size=5
## + : mtry=5, splitrule=gini, min.node.size=5
## - : mtry=5, splitrule=gini, min.node.size=5
## + : mtry=6, splitrule=gini, min.node.size=5
## - : mtry=6, splitrule=gini, min.node.size=5
## Aggregating results
## Selecting tuning parameters
## Fitting mtry = 5, splitrule = gini, min.node.size = 5 on full training set
```

```
randomforest_over
```

```
## Random Forest
##
## 4543 samples
## 10 predictor
## 2 classes: '1', '2'
##
## No pre-processing
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.9511336 0.9021617
## 3 0.9575171 0.9149570
```

```
## 4      0.9610390  0.9220165
## 5      0.9630200  0.9259846
## 6      0.9603786  0.9207001
##
## Tuning parameter 'splitrule' was held constant at a value of gini
##
## Tuning parameter 'min.node.size' was held constant at a value of 5
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were mtry = 5, splitrule = gini
## and min.node.size = 5.

# Making predictions based off of testing data
pred_randomforest_over<- predict(randomforest_over, newdata=beta_testing_over)

# Create a confusion matrix
confusionMatrix(pred_randomforest_over, factor(testing_over[["stroke"]]), positive = "1")

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  1    2
##           1 917  44
##           2  18 968
##
##               Accuracy : 0.9682
##               95% CI : (0.9594, 0.9755)
##       No Information Rate : 0.5198
##       P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.9363
##
##  Mcnemar's Test P-Value : 0.001498
##
##       Sensitivity : 0.9807
##       Specificity : 0.9565
##       Pos Pred Value : 0.9542
##       Neg Pred Value : 0.9817
##       Prevalence : 0.4802
##       Detection Rate : 0.4710
##       Detection Prevalence : 0.4936
##       Balanced Accuracy : 0.9686
##
##       'Positive' Class : 1
##
```

Model 6: Extreme gradient boosting tree

```
# Build the model using the training data
xgb_over <- train(
  stroke ~ ., training_over, method = "xgbTree", tuneLength = 3, tuneGrid = xgbgrid, trControl = xgbcon
)
```

```
xgb_over
```

```
## eXtreme Gradient Boosting
##
## 4543 samples
## 10 predictor
## 2 classes: '1', '2'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 3635, 3634, 3634, 3635, 3634
## Resampling results:
##
## Accuracy Kappa
## 0.9621395 0.9242213
##
## Tuning parameter 'nrounds' was held constant at a value of 3500
##
## Tuning parameter 'min_child_weight' was held constant at a value of 0
##
## Tuning parameter 'subsample' was held constant at a value of 0.5
```

```
# Making predictions based off of testing data
pred_xgb_over <- predict(xgb_over, newdata = beta_testing_over)

# Create a confusion matrix
confusionMatrix(pred_xgb_over, factor(testing_over[["stroke"]]), positive = "1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  1    2
##           1 921  47
##           2  14 965
##
##           Accuracy : 0.9687
##           95% CI : (0.9599, 0.976)
##           No Information Rate : 0.5198
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9373
##
## Mcnemar's Test P-Value : 4.182e-05
##
##           Sensitivity : 0.9850
##           Specificity : 0.9536
##           Pos Pred Value : 0.9514
##           Neg Pred Value : 0.9857
##           Prevalence : 0.4802
##           Detection Rate : 0.4730
##           Detection Prevalence : 0.4972
##           Balanced Accuracy : 0.9693
```

```
##  
##      'Positive' Class : 1  
##
```